

# Error handling

# Objectives

- Explore Go error management
- Understand:
  - Errors
  - panic()
  - recover()
  - defer

# Defer

- **defer statement**
  - An ordinary function call prefixed by the defer keyword
- Deferred function and **argument expressions are evaluated when the statement is executed**
- The actual **call is deferred** until the function completes
  - By return statement, falling off the end, exiting or panicking
- **Any number of calls may be deferred**
  - Deferred calls are executed in **reverse order** in which they were deferred
- A defer statement is often used with paired operations
  - e.g. open/close, connect/disconnect, lock/unlock
  - Performs like a finally block in some languages, ensuring that resources are released in all cases
  - The **right place for a defer statement is immediately after the resource has been successfully acquired**
- Deferred functions run after return statements have updated the function's result variables
  - An anonymous function can access its enclosing function's variables, including named results, even change them ... (!)
- The example program defers the HTTP response body close operation
  - Also demonstrates use of nested anonymous functions for internal recursion
  - This function uses a closure to call itself (required since the function is anonymous)
  - The code uses the golang net/html extension package (not part of the std lib) to parse the HTML document

```
1 package main
2
3 import (
4     "fmt"
5     "net/http"
6     "strings"
7
8     "golang.org/x/net/html"
9 )
10
11 func title(url string) (*html.Node, error) {
12     resp, err := http.Get(url)
13     if err != nil {
14         return nil, err
15     }
16     defer resp.Body.Close()
17
18     ct := resp.Header.Get("Content-Type")
19     if ct != "text/html" && !strings.HasPrefix(ct, "text/html;") {
20         return nil, fmt.Errorf("%s has type %s, not text/html", url, ct)
21     }
22     doc, err := html.Parse(resp.Body)
23     if err != nil {
24         return nil, fmt.Errorf("parsing %s as HTML: %v", url, err)
25     }
26     var f func(n *html.Node) *html.Node
27     f = func(n *html.Node) *html.Node {
28         if n.Type == html.ElementNode && n.Data == "title" {
29             return n.FirstChild
30         }
31         for c := n.FirstChild; c != nil; c = c.NextSibling {
32             p := f(c)
33             if p != nil {
34                 return p
35             }
36         }
37         return nil
38     }
39     return f(doc), nil
40 }
41
42 func main() {
43     el, _ := title("http://rx-m.com")
44     fmt.Println(el.Data)
45 }
46
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

```
2017/04/10 20:51:43 server.go:73: Using API v1
2017/04/10 20:51:43 debugger.go:68: launching process with args: [d:\dev\go\exampl
API server listening at: 127.0.0.1:2345
2017/04/10 20:51:43 debugger.go:414: continuing
RX-M - DevOps & Cloud Native Consulting and Training - RX-M
```

# Errors

4

Copyright 2013-2017, RX-M LLC

- Some functions **always succeed** (200)
- Some functions **fail if preconditions are not met** (400)
  - Errors are a result of bad calls
- Some functions given proper preconditions **fail due to factors beyond the functions control** (500)
  - Any function that does I/O
- Errors are an important part of a package's API
  - Go considers errors one of a set of expected behaviors
  - If a function can fail it is expected to return an error result
    - Conventionally the last thing in the return list
  - If the failure has only one possible cause the result is (by convention) a Boolean called "ok"

```
1 package main
2
3 import "fmt"
4
5 func div(x, y int) (res int, ok bool) {
6     if y == 0 {
7         return 0, false
8     }
9     return x / y, true
10 }
11
12 func main() {
13     ans, ok := div(8, 4)
14     if ok {
15         fmt.Println("First try: ", ans)
16     }
17     ans, ok = div(8, 0)
18     if ok {
19         fmt.Println("Second try: ", ans)
20     }
21 }
22
```

PROBLEMS	OUTPUT	DEBUG CONSOLE	TERMINAL
			2017/04/10 15:11:57 server.go:73: Using API v1
			2017/04/10 15:11:57 debugger.go:68: launching pro
			API server listening at: 127.0.0.1:2345
			2017/04/10 15:11:57 debugger.go:414: continuing
			First try: 2

# error

5

Copyright 2013-2017, RX-M LLC

- Non trivial (bool) errors are indicated by returning an “error” instance as a final return value
- **error is a built-in interface type**
  - An error may be:
    - **nil** - implies success
    - **non-nil** - implies failure
      - Such an error has a message string obtained by calling its Error method
      - To print the message: `fmt.Println(err)` or `fmt.Printf("%v", err)` or `fmt.Println(err.Error())`
      - Usually when a function returns a non-nil error its other results are undefined
        - Some rare functions return partial results in error cases
- **Go does not offer a traditional exception handling system**
  - Go supports a type of exception for reporting truly unexpected errors that indicate a bug (not to be used for reporting routine errors)
  - The Go philosophy is that exceptions entangle the description of an error with the control flow required to handle it
    - This leads to routine errors being reported to the end user in a stack trace, full of information about the structure of the program but lacking intelligible context about what went wrong
  - **Go programs use ordinary control-flow mechanisms (e.g. if and return) to respond to errors**
    - This style undeniably demands that more attention be paid to error-handling logic
    - precisely the point

```
1 package main
2
3 import (
4     "errors"
5     "fmt"
6 )
7
8 func posdiv(x, y int) (res int, err error) {
9     if y < 0 {
10         return 0, errors.New("attempted posdiv division with a negative number")
11     }
12     if y == 0 {
13         return 0, errors.New("attempted division by 0")
14     }
15     return x / y, nil
16 }
17
18 func main() {
19     ans, err := posdiv(8, -4)
20     if err != nil {
21         fmt.Println("Error: ", err.Error())
22     } else {
23         fmt.Println("Result: ", ans)
24     }
25 }
26
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

```
2017/04/10 15:57:03 server.go:73: Using API v1
2017/04/10 15:57:03 debugger.go:68: launching process with args: [d:\dev\go\example\src\
API server listening at: 127.0.0.1:2345
2017/04/10 15:57:03 debugger.go:414: continuing
Error: attempted posdiv division with a negative number
```

# Propagating errors

6

Copyright 2013-2017, RX-M LLC

- Errors in low level routines must often be propagated by mid level routines that have no stake in the error
  - return err
- A call which produces an error such that the callee does not have enough context to produce a meaningful explanation may require the caller to augment it to keep the error meaningful
  - Go allows new errors to be fabricated from lower level error messages
    - fmt.Errorf("%v", err)
    - errors.Wrap(err, "read failed")
  - Error messages are frequently chained in this way
- Therefore error messages should:
  - Avoid newlines
  - Not be capitalized
  - Provide a meaningful description of the problem
  - Supply sufficient/relevant detail
  - Be consistent
    - errors returned by the same function/package should be of similar form
    - e.g. The os package guarantees that every error returned by a file operation describes the nature of the failure (permission denied) and the name of the file
- Resulting errors may be long, but they will be self-contained and easily parsed by tools like grep

```
1 package main
2
3 import "errors"
4 import "fmt"
5
6 func posdiv(x, y int) (res int, err error) {
7     if y < 0 {
8         return 0, errors.New("attempted posdiv division with a negative number")
9     }
10    if y == 0 {
11        return 0, errors.New("attempted division by 0")
12    }
13    return x / y, nil
14 }
15
16 func multidiv(x, y1, y2 int) (res int, err error) {
17     ans1, e1 := posdiv(x, y1)
18     ans2, e2 := posdiv(x, y2)
19     if e1 != nil || e2 != nil {
20         return 0, fmt.Errorf("multidiv component failure, y1: %v, y2: %v", e1, e2)
21     }
22     return ans1 + ans2, nil
23 }
24
25 func main() {
26     ans, err := multidiv(8, 2, -4)
27     if err != nil {
28         fmt.Println("Error: ", err.Error())
29     } else {
30         fmt.Println("Result: ", ans)
31     }
32 }
33
```

PROBLEMS	OUTPUT	DEBUG CONSOLE	TERMINAL
			2017/04/10 16:09:56 server.go:73: Using API v1
			2017/04/10 16:09:56 debugger.go:68: launching process with args: [d:\dev\go\example\src\debug]
			API server listening at: 127.0.0.1:2345
			2017/04/10 16:09:56 debugger.go:414: continuing
			Error: multidiv component failure, y1: <nil>, y2: attempted posdiv division with a negative number

# Retry Errors

7

Copyright 2013-2017, RX-M LLC

- Transient/unpredictable problems may require code to retry the failed operation
  - Perhaps with a delay
  - Perhaps with a limit on the number of attempts or time spent
  - Transient errors recovered from should not go unrecorded
- The Go **log** package
  - Implements simple logging
  - Defines a **Logger** type with methods for formatting output
  - Has a predefined 'standard' **Logger** accessible through helper functions **Print[f|ln]**, **Fatal[f|ln]**, and **Panic[f|ln]**
    - The **Fatal** functions call **os.Exit(1)** after writing the log message
    - The **Panic** functions call **panic** after writing the log message
  - The standard logger writes to standard error and prints the date and time of each logged message
  - Every log message is output on a separate line
    - if the message being printed does not end in a newline, the logger will add one
- **Stderr** can be printed to directly with the **fmt** package
  - **fmt.Fprintf(os.Stderr, "big problem!")**

```
1 package main
2
3 import "log"
4
5 func main() {
6     log.Printf("connection failed retrying")
7 }
8
```

PROBLEMS	OUTPUT	DEBUG CONSOLE	TERMINAL
	2017/04/10 16:31:22	server.go:73: Using API v1	
	2017/04/10 16:31:22	debugger.go:68: launching process	
		API server listening at: 127.0.0.1:2345	
	2017/04/10 16:31:23	debugger.go:414: continuing	
	2017/04/10 16:31:23	connection failed retrying	

# Stopping Execution

- If the error is unrecoverable it may be best to terminate execution
  - The os package provides an `func Exit(code int)` function to terminate execution
    - The code passed to Exit is passed to the system as the process exit code (0 typically indicates success, all other values indicate failure)
  - This decision should be reserved for the most abstract parts of the program (e.g. `main()`)
- Go error handling style
  - Invoke risky function
    - Checking for error
    - Deal with Failure
      - If failure causes the function to return the success logic follows at the outer level
    - Deal with Success
  - Functions include `initial checks to reject errors` (DbC\*) followed by the substance of the function minimally indented

*\* In Design by Contract style*



# Distinguished Errors

- Some packages define specific errors
  - Referred to as distinguished errors
  - `io.EOF` is a good example
    - In an end-of-file condition there is no information to report
    - This error has a fixed message: "EOF"
  - **if `err == io.EOF`**
    - A legal expression because the `io.EOF` variable is always returned when an EOF event occurs

EOF is the error returned by `Read` when no more input is available. Functions should return EOF only to signal a graceful end of input. If the EOF occurs unexpectedly in a structured data stream, the appropriate error is either `ErrUnexpectedEOF` or some other error giving more detail.

```
var EOF = errors.New("EOF")
```

io package errors

`ErrClosedPipe` is the error used for read or write operations on a closed pipe.

```
var ErrClosedPipe = errors.New("io: read/write on closed pipe")
```

`ErrNoProgress` is returned by some clients of an `io.Reader` when many calls to `Read` have failed to return any data or error, usually the sign of a broken `io.Reader` implementation.

```
var ErrNoProgress = errors.New("multiple Read calls return no data or error")
```

`ErrShortBuffer` means that a read required a longer buffer than was provided.

```
var ErrShortBuffer = errors.New("short buffer")
```

`ErrShortWrite` means that a write accepted fewer bytes than requested but failed to return an explicit error.

```
var ErrShortWrite = errors.New("short write")
```

`ErrUnexpectedEOF` means that EOF was encountered in the middle of reading a fixed-size block or data structure.

```
var ErrUnexpectedEOF = errors.New("unexpected EOF")
```

# Panics

10

Copyright 2013-2017, RX-M LLC

- In Go execution errors such as attempting to index an array out of bounds trigger a run-time panic
  - Can also be triggered by the built-in function `panic(string)`
- Panics are designed to fail fast, stopping all operations quickly
- A common use of panic is to abort when a function returns an error value that we don't know how to (or want to) handle
  - Last resort, idiomatic go code returns meaningful errors
- During a typical panic, **normal execution stops**, all **deferred function calls are executed**, and the **program crashes with a log message**
  - This log message includes the panic value, which is usually an error message of some sort a stack trace showing the stack of function calls that were active at the time of the panic
  - This log message often has enough information to diagnose the root cause of the problem without running the program again, so it should always be included in a bug report about a panicking program

```
1 package main
2
3 func main() {
4     panic(3)
5 }
6
```

PROBLEMS	OUTPUT	DEBUG CONSOLE	TERMINAL
	2017/04/10 21:13:04	server.go:73: Using API v1	
	2017/04/10 21:13:04	debugger.go:68: launching proc	
		API server listening at: 127.0.0.1:2345	
	2017/04/10 21:13:04	debugger.go:414: continuing	
	2017/04/10 21:13:08	debugger.go:414: continuing	
		panic: 3	
		goroutine 1 [running]:	
		main.main()	
		d:/dev/go/example/src/example.go:4 +0x5f	

# Recover

11

Copyright 2013-2017, RX-M LLC

- **Giving up** is usually the right response to a panic
- In some cases **it may be possible to recover**
  - Or clean up before quitting (e.g. a web server could close connections before crashing)
- If the built-in **recover()** function is called within a deferred function that is panicking the current state of panic ends and recover() returns the panic code
  - The function that was panicking does not continue where it left off but returns normally
- If recover is called at any other time, it has no effect and returns nil
- The example defers an anonymous function (self invoking due to the trailing parenthesis)
  - The panic code was 3 so the Println() displays this as the result of recover()
  - The program exits normally

```
1 package main
2
3 import "fmt"
4
5 func main() {
6     defer func() {
7         fmt.Println(recover())
8     }()
9     panic(3)
10 }
11
```

PROBLEMS	OUTPUT	DEBUG CONSOLE	TERMINAL
	2017/04/10 21:22:38	server.go:73: Using API v1	
	2017/04/10 21:22:38	debugger.go:68: launching process	
		API server listening at: 127.0.0.1:2345	
	2017/04/10 21:22:38	debugger.go:414: continuing	
		3	

# Summary

- Go provides unique Go specific function features
  - Deferred functions, no exceptions, ...
- Go error management is explicit using tools such as:
  - Errors
  - panic()
  - recover()
  - defer

# Lab: Error handling

- Error handling